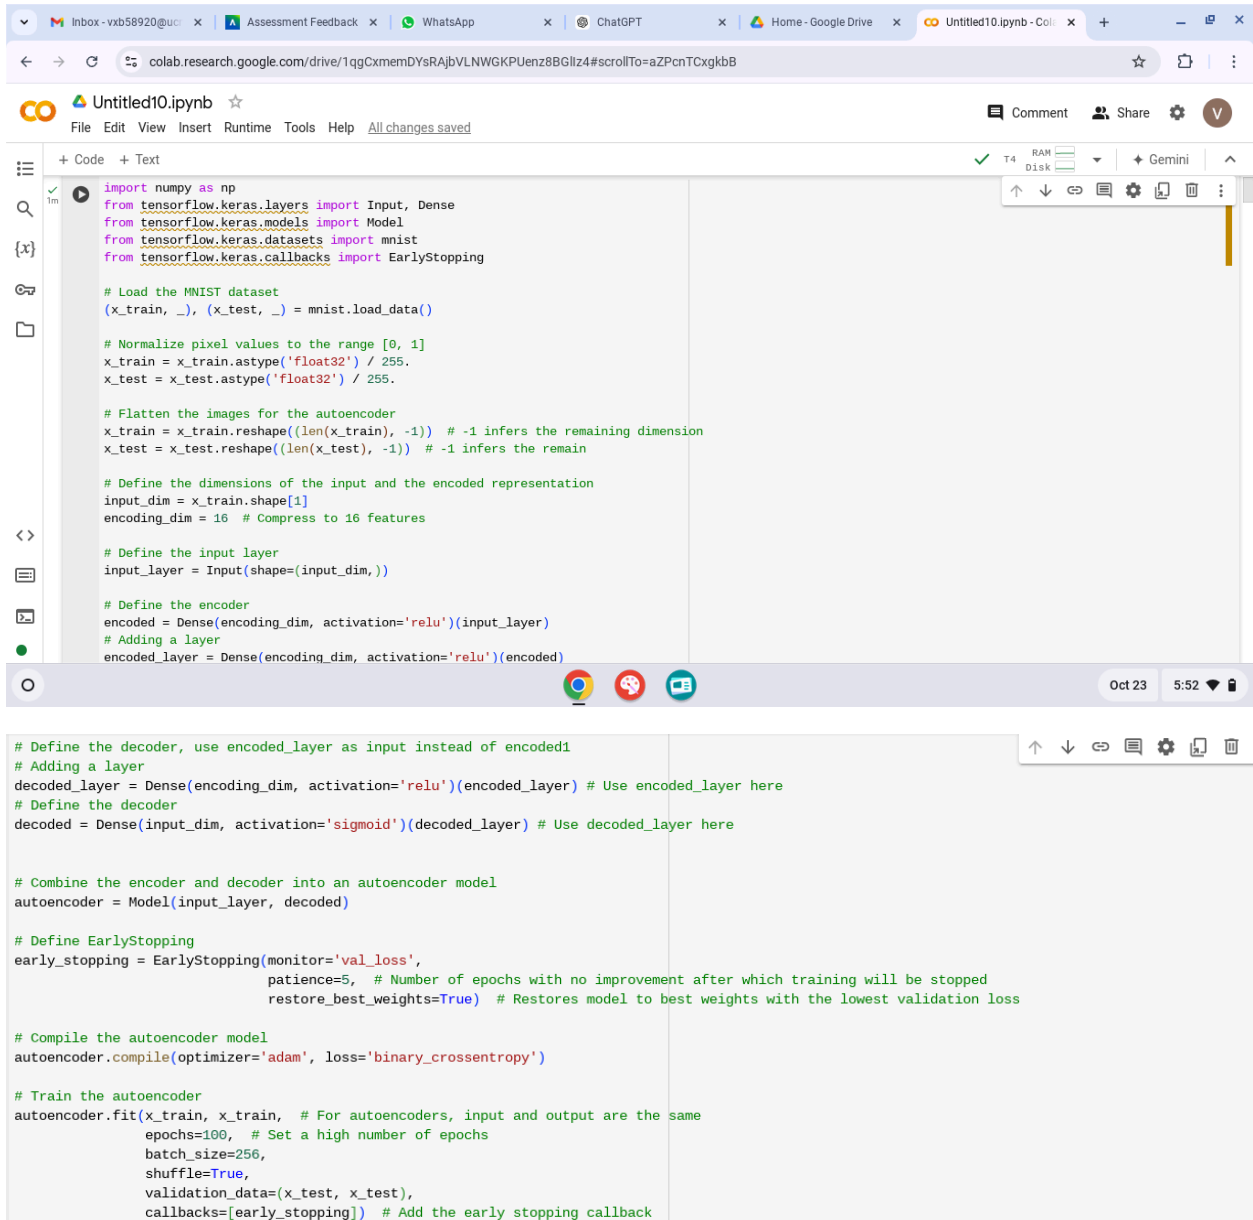


ICP 6 REPORT



```
import numpy as np
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
from tensorflow.keras.datasets import mnist
from tensorflow.keras.callbacks import EarlyStopping

# Load the MNIST dataset
(x_train, _), (x_test, _) = mnist.load_data()

# Normalize pixel values to the range [0, 1]
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.

# Flatten the images for the autoencoder
x_train = x_train.reshape((len(x_train), -1)) # -1 infers the remaining dimension
x_test = x_test.reshape((len(x_test), -1)) # -1 infers the remain

# Define the dimensions of the input and the encoded representation
input_dim = x_train.shape[1]
encoding_dim = 16 # Compress to 16 features

# Define the input layer
input_layer = Input(shape=(input_dim,))

# Define the encoder
encoded = Dense(encoding_dim, activation='relu')(input_layer)
# Adding a layer
encoded_layer = Dense(encoding_dim, activation='relu')(encoded)

# Define the decoder, use encoded_layer as input instead of encoded1
# Adding a layer
decoded_layer = Dense(encoding_dim, activation='relu')(encoded_layer) # Use encoded_layer here
# Define the decoder
decoded = Dense(input_dim, activation='sigmoid')(decoded_layer) # Use decoded_layer here

# Combine the encoder and decoder into an autoencoder model
autoencoder = Model(input_layer, decoded)

# Define EarlyStopping
early_stopping = EarlyStopping(monitor='val_loss',
                               patience=5, # Number of epochs with no improvement after which training will be stopped
                               restore_best_weights=True) # Restores model to best weights with the lowest validation loss

# Compile the autoencoder model
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

# Train the autoencoder
autoencoder.fit(x_train, x_train, # For autoencoders, input and output are the same
               epochs=100, # Set a high number of epochs
               batch_size=256,
               shuffle=True,
               validation_data=(x_test, x_test),
               callbacks=[early_stopping]) # Add the early stopping callback
```

```

import numpy as np
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
from tensorflow.keras.datasets import mnist
from tensorflow.keras.callbacks import TerminateOnNaN

# Define the TerminateOnNaN callback
terminate_on_nan = TerminateOnNaN()

# Load the MNIST dataset
(x_train, _), (x_test, _) = mnist.load_data()

# Normalize pixel values to the range [0, 1]
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.

# Flatten the images for the autoencoder
x_train = x_train.reshape((len(x_train), -1)) # -1 infers the remaining dimension
x_test = x_test.reshape((len(x_test), -1)) # -1 infers the remain

# Define the dimensions of the input and the encoded representation
input_dim = x_train.shape[1]
encoding_dim = 16 # Compress to 16 features

# Define the input layer
input_layer = Input(shape=(input_dim,))

```

```

# Define the encoder
encoded = Dense(encoding_dim, activation='relu')(input_layer)
# Adding a layer
encoded_layer = Dense(encoding_dim, activation='relu')(encoded)

# Adding a layer, use encoded_layer as input instead of encoded1
decoded_layer = Dense(encoding_dim, activation='relu')(encoded_layer) # Use encoded_layer here
# Define the decoder, use decoded_layer as input instead of decoded1
decoded = Dense(input_dim, activation='sigmoid')(decoded_layer) # Use decoded_layer here

# Combine the encoder and decoder into an autoencoder model
autoencoder = Model(input_layer, decoded)

# Compile the autoencoder model
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

# Train the autoencoder
# Assuming x_train and x_test are your training and validation datasets
autoencoder.fit(x_train, x_train, # For autoencoders, input and output are the same
               epochs=30, # Set the number of epochs
               batch_size=256,
               shuffle=True,
               validation_data=(x_test, x_test),
               callbacks=[terminate_on_nan]) # Add the TerminateOn

```



```

import numpy as np
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
from tensorflow.keras.datasets import mnist
from tensorflow.keras.callbacks import ReduceLRonPlateau

# Define the ReduceLRonPlateau callback
reduce_lr = ReduceLRonPlateau(monitor='val_loss', # Metric to monitor
                             factor=0.5, # Factor by which the learning rate will be reduced (new_lr = lr * factor)
                             patience=3, # Number of epochs with no improvement after which learning rate will be reduced
                             min_lr=1e-6, # Lower bound for the learning rate
                             verbose=1) # Print message when the learning rate is reduced

# Load the MNIST dataset
(x_train, _), (x_test, _) = mnist.load_data()

# Normalize pixel values to the range [0, 1]
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.

# Flatten the images for the autoencoder
x_train = x_train.reshape((len(x_train), -1)) # -1 infers the remaining dimension
x_test = x_test.reshape((len(x_test), -1)) # -1 infers the remain

# Define the dimensions of the input and the encoded representation
input_dim = x_train.shape[1]
encoding_dim = 16 # Compress to 16 features

```

```

# Define the input layer
input_layer = Input(shape=(input_dim,))

# Define the encoder
encoded = Dense(encoding_dim, activation='relu')(input_layer)
# Adding a layer
encoded_layer = Dense(encoding_dim, activation='relu')(encoded)

# Adding a layer, use encoded_layer as input instead of encoded
decoded_layer = Dense(encoding_dim, activation='relu')(encoded_layer) # Use encoded_layer here
# Define the decoder, use decoded_layer as input instead of decoded
decoded = Dense(input_dim, activation='sigmoid')(decoded_layer) # Use decoded_layer here

# Combine the encoder and decoder into an autoencoder model
autoencoder = Model(input_layer, decoded)

# Compile the autoencoder model
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

# Train the autoencoder
# Assuming x_train and x_test are your training and validation datasets
autoencoder.fit(x_train, x_train, # For autoencoders, input and output are the same
               epochs=30, # Number of epochs
               batch_size=256,
               shuffle=True,
               validation_data=(x_test, x_test), # Validation data
               callbacks=[reduce_lr]) # Add the ReduceLRonPlateau callback

```

```

import numpy as np
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
from tensorflow.keras.datasets import mnist
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint, TerminateOnNaN, ReduceLRonPlateau

# EarlyStopping callback to stop training if validation loss stops improving
early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)

# ModelCheckpoint callback to save the best model based on validation loss
checkpoint = ModelCheckpoint(filepath='autoencoder_best.keras', monitor='val_loss', save_best_only=True, verbose=1)

# TerminateOnNaN callback to stop training if the loss becomes NaN
terminate_on_nan = TerminateOnNaN()

# Define the ReduceLRonPlateau callback
reduce_lr = ReduceLRonPlateau(monitor='val_loss', factor=0.5, patience=3, min_lr=1e-6, verbose=1)

# Load the MNIST dataset
(x_train, _), (x_test, _) = mnist.load_data()

# Normalize pixel values to the range [0, 1]
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.

# Flatten the images for the autoencoder
x_train = x_train.reshape((len(x_train), -1)) # -1 infers the remaining dimension
x_test = x_test.reshape((len(x_test), -1)) # -1 infers the remain

```

```
# Define the dimensions of the input and the encoded representation
input_dim = x_train.shape[1]
encoding_dim = 16 # Compress to 16 features

# Define the input layer
input_layer = Input(shape=(input_dim,))

# Define the encoder
encoded = Dense(encoding_dim, activation='relu')(input_layer)
# Adding a layer
encoded_layer = Dense(encoding_dim, activation='relu')(encoded)

# Adding a layer, use encoded_layer as input instead of encoded1
decoded_layer = Dense(encoding_dim, activation='relu')(encoded_layer) # Fixed: Use encoded_layer here
# Define the decoder, use decoded_layer as input instead of decoded1
decoded = Dense(input_dim, activation='sigmoid')(decoded_layer) # Fixed: Use decoded_layer here

# Combine the encoder and decoder into an autoencoder model
autoencoder = Model(input_layer, decoded)

# Compile the autoencoder model
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

# Training with multiple callbacks
autoencoder.fit(x_train, x_train,
               epochs=30, # You can set a high number of epochs
               batch_size=256,
               shuffle=True,
               validation_data=(x_test, x_test),
               callbacks=[reduce_lr, early_stopping, checkpoint, terminate_on_nan]) # Using multiple callbacks
```

```
Epoch 17/30
228/235 ----- 0s 2ms/step - loss: 0.1469
Epoch 17: val_loss improved from 0.14541 to 0.14494, saving model to autoencoder_best.keras
235/235 ----- 1s 3ms/step - loss: 0.1469 - val_loss: 0.1449 - learning_rate: 0.0010
Epoch 18/30
229/235 ----- 0s 2ms/step - loss: 0.1465
Epoch 18: val_loss improved from 0.14494 to 0.14440, saving model to autoencoder_best.keras
235/235 ----- 1s 3ms/step - loss: 0.1465 - val_loss: 0.1444 - learning_rate: 0.0010
Epoch 19/30
231/235 ----- 0s 2ms/step - loss: 0.1461
Epoch 19: val_loss improved from 0.14440 to 0.14390, saving model to autoencoder_best.keras
235/235 ----- 1s 3ms/step - loss: 0.1461 - val_loss: 0.1439 - learning_rate: 0.0010
Epoch 20/30
225/235 ----- 0s 2ms/step - loss: 0.1457
Epoch 20: val_loss improved from 0.14390 to 0.14365, saving model to autoencoder_best.keras
235/235 ----- 1s 3ms/step - loss: 0.1457 - val_loss: 0.1437 - learning_rate: 0.0010
Epoch 21/30
228/235 ----- 0s 2ms/step - loss: 0.1446
Epoch 21: val_loss improved from 0.14365 to 0.14300, saving model to autoencoder_best.keras
235/235 ----- 1s 3ms/step - loss: 0.1446 - val_loss: 0.1430 - learning_rate: 0.0010
Epoch 22/30
226/235 ----- 0s 2ms/step - loss: 0.1443
Epoch 22: val_loss improved from 0.14300 to 0.14253, saving model to autoencoder_best.keras
```

```
[14] from tensorflow.keras.models import load_model

# Load the entire model
best_autoencoder = load_model('autoencoder_best.keras')

# Let's look at the encoded representations
encoded_data = best_autoencoder.predict(x_test)
print(encoded_data)
print(encoded_data.shape)
```

```
313/313 ----- 1s 3ms/step
[[[1.75119294e-11 5.90835920e-13 3.34731982e-12 ... 9.35920265e-12
 1.14782457e-12 5.60907223e-13]
 [3.73498033e-09 4.47984192e-08 2.67801195e-08 ... 1.48025379e-08
 7.60447794e-09 1.10647704e-07]
 [9.43129949e-11 2.13933468e-10 9.09995745e-10 ... 3.27405131e-10
 3.95407387e-11 2.15216733e-08]
 ...
 [1.53404189e-15 1.70396599e-16 3.91804577e-15 ... 2.86246039e-15
 3.23801558e-16 2.59001584e-15]
 [1.27121022e-11 9.90096980e-12 8.65826369e-11 ... 4.37883167e-11
 6.05044834e-12 6.50171861e-11]
 [3.22090504e-15 1.58808479e-14 6.83981053e-15 ... 2.27953473e-14
 5.10732788e-15 2.07234864e-13]]
(10000, 784)]
```

My Github Link :

<https://github.com/Nitish300903/bda.git>